


Lab 5 Demo: Packet Sniffing and Spoofing

Stephanie Salgado

Set up:



Located lab setup files in shared folder.

```
stephaniesalgado@VM: ~/Share/Labsetup
stephaniesalgado@VM:~/Share/Labsetup$ sudo docker-compose build
attacker uses an image, skipping
hostA uses an image, skipping
hostB uses an image, skipping
stephaniesalgado@VM:~/Share/Labsetup$ sudo docker-compose up
Creating network "net-10.9.0.0" with the default driver
Creating hostA-10.9.0.5 ... done
Creating seed-attacker ... done
Creating hostB-10.9.0.6 ... done
Attaching to seed-attacker, hostA-10.9.0.5, hostB-10.9.0.6
hostA-10.9.0.5 | * Starting internet superserver inetd      [ OK ]
hostB-10.9.0.6 | * Starting internet superserver inetd      [ OK ]
```

I began by navigating to project files and using “sudo docker-compose build” then “sudo docker-compose up” to build and start the container.

```

stephaniesalgado@VM: ~/Share/Labsetup$ sudo docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS   NAMES
f589e678e850   handsonsecurity/seed-ubuntu:large   "bash -c ' /etc/init..." 16 minutes ago Up 16 minutes        hostB-10.9.0.6
d658205a3f98   handsonsecurity/seed-ubuntu:large   "/bin/sh -c /bin/bash"    16 minutes ago Up 16 minutes        seed-attacker
2364903ac6b4   handsonsecurity/seed-ubuntu:large   "bash -c ' /etc/init..." 16 minutes ago Up 16 minutes        hostA-10.9.0.5
stephaniesalgado@VM: ~/Share/Labsetup$ sudo docker exec -it seed-attacker bash
root@VM: /#

```

In another terminal, I used “sudo docker ps” to find out the ID of the container. Then, “sudo docker exec -it seed-attacker bash” to start a shell on the container.

```

root@VM: /# ifconfig
br-05ff62bc2439: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 10.100.0.1 netmask 255.255.255.0 broadcast 10.100.0.255
    ether 02:42:d3:e3:a3:a0 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-0ebf74fb72d4: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.23.0.1 netmask 255.255.0.0 broadcast 172.23.255.255
    ether 02:42:f2:ee:0b:07 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-204b65107357: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 10.151.0.1 netmask 255.255.255.0 broadcast 10.151.0.255
    ether 02:42:c9:36:0c:a1 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-3a4b4583bca0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 10.153.0.1 netmask 255.255.255.0 broadcast 10.153.0.255
    ether 02:42:f5:7d:53:95 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-9a8a3d9d41a2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.1 netmask 255.255.255.0 broadcast 10.9.0.255
    inet6 fe80::42:94ff:fe6c:f61d prefixlen 64 scopeid 0x20<link>
    ether 02:42:94:0c:f6:1d txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 49 bytes 6026 (6.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-aeac8ab296fe: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 10.152.0.1 netmask 255.255.255.0 broadcast 10.152.0.255
    ether 02:42:7d:05:23:d1 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)

```

I used “ifconfig” to find the name of the corresponding network interface on my VM. I made sure to look for IP address 10.9.0.1. That’s where I got, “br-9a8a3d9d41a2”.

```

stephaniesalgado@VM: ~/Share/Labsetup$ sudo docker network ls
NETWORK ID      NAME                                DRIVER  SCOPE
2241646630cb    bridge                             bridge  local
b3581338a28d    host                               host    local
f438ea825ed2    internet-nano_default              bridge  local
204b65107357    internet-nano_net_151_net0         bridge  local
aeac8ab296fe    internet-nano_net_152_net0         bridge  local
3a4b4583bca0    internet-nano_net_153_net0         bridge  local
05ff62bc2439    internet-nano_net_ix_ix100         bridge  local
0ebf74fb72d4    map_default                         bridge  local
9a8a3d9d41a2    net-10.9.0.0                       bridge  local
77aceccbe26     none                               null    local

```

I made sure by using the “sudo docker network ls” command.

Task 1.1: Sniffing Packets

```
stephaniesalgado@VM: ~/Share/Labsetup

root@VM:/volumes# nano task1.1.py
root@VM:/volumes# cat task1.1.py
#!/usr/bin/env python3
from scapy.all import *

print("SNIFFING PACKETS.....")

def print_pkt(pkt):
    pkt.show()

pkt = sniff(iface='br-9a8a3d9d41a2', filter='icmp', prn=print_pkt)
root@VM:/volumes# chmod a+x task1.1.py
root@VM:/volumes# ./task1.1.py
SNIFFING PACKETS.....
```

I navigated into “volumes” on the attack container then created the “task1.1.py” file. After that, using “chmod a+x task1.1.py” I made the file executable. Since I’m already “root” by default in the attack container, I only had to use “./task1.1.py” to run the file.

```
stephaniesalgado@VM:~/Share/Labsetup$ sudo docker exec -it hostA-10.9.0.5 bash
root@2364903ac6b4:/# ping 10.9.0.6
PING 10.9.0.6 (10.9.0.6) 56(84) bytes of data.
64 bytes from 10.9.0.6: icmp_seq=1 ttl=64 time=0.139 ms
64 bytes from 10.9.0.6: icmp_seq=2 ttl=64 time=0.050 ms
64 bytes from 10.9.0.6: icmp_seq=3 ttl=64 time=0.062 ms
64 bytes from 10.9.0.6: icmp_seq=4 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=5 ttl=64 time=0.052 ms
64 bytes from 10.9.0.6: icmp_seq=6 ttl=64 time=0.052 ms
64 bytes from 10.9.0.6: icmp_seq=7 ttl=64 time=0.064 ms
```

I created network traffic by going to “hostA” and pinging “hostB” with command “ping 10.9.0.6”, where “10.9.0.6” is the IP corresponding to “hostB”. I let that run for a while.

```
stephaniesalgado@VM: ~/Share/Labsetup

64 bytes from 10.9.0.6: icmp_seq=85 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=86 ttl=64 time=0.067 ms
64 bytes from 10.9.0.6: icmp_seq=87 ttl=64 time=0.046 ms
64 bytes from 10.9.0.6: icmp_seq=88 ttl=64 time=0.044 ms
64 bytes from 10.9.0.6: icmp_seq=89 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=90 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=91 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=92 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=93 ttl=64 time=0.093 ms
64 bytes from 10.9.0.6: icmp_seq=94 ttl=64 time=0.064 ms
64 bytes from 10.9.0.6: icmp_seq=95 ttl=64 time=0.053 ms
64 bytes from 10.9.0.6: icmp_seq=96 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=97 ttl=64 time=0.053 ms
64 bytes from 10.9.0.6: icmp_seq=98 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=99 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=100 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=101 ttl=64 time=0.060 ms
64 bytes from 10.9.0.6: icmp_seq=102 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=103 ttl=64 time=0.048 ms
64 bytes from 10.9.0.6: icmp_seq=104 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=105 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=106 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=107 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=108 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=109 ttl=64 time=0.043 ms
64 bytes from 10.9.0.6: icmp_seq=110 ttl=64 time=0.048 ms
64 bytes from 10.9.0.6: icmp_seq=111 ttl=64 time=0.045 ms
64 bytes from 10.9.0.6: icmp_seq=112 ttl=64 time=0.048 ms
64 bytes from 10.9.0.6: icmp_seq=113 ttl=64 time=0.045 ms
^C
--- 10.9.0.6 ping statistics ---
113 packets transmitted, 113 received, 0% packet loss, time 11474ms
rtt min/avg/max/mdev = 0.041/0.055/0.139/0.016 ms
root@2364903ac6b4:/#
```

I stopped after 113 packets.


```
stephaniesalgado@VM: ~/Share/Labsetup
root@VM:/volumes# nano task1.1.py
root@VM:/volumes# ./task1.1.py
SNIFFING PACKETS.....

===== PACKET: 1 =====

###[ Ethernet ]###
dst      = 02:42:0a:09:00:05
src      = 02:42:0a:09:00:06
type     = IPv4
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x10
len      = 60
id       = 16684
flags    = DF
frag     = 0
ttl      = 64
proto    = tcp
chksum   = 0xe563
src      = 10.9.0.6
dst      = 10.9.0.5
\options \
###[ TCP ]###
sport    = 55112
dport    = telnet
seq      = 1947937917
ack      = 0
dataoffs = 10
reserved = 0
flags    = S
window   = 64240
chksum   = 0x144b
urgptr   = 0
options  = [('MSS', 1460), ('SackOK', b'), ('Timestamp', (2714716888, 0)), ('NOP', None), ('WScale', 7)]

===== PACKET: 2 =====

###[ Ethernet ]###
dst      = 02:42:0a:09:00:05
src      = 02:42:0a:09:00:06
type     = IPv4
```

This is the output I got after filtering with “filter=tcp && src host 10.9.06 && dst port 23”.

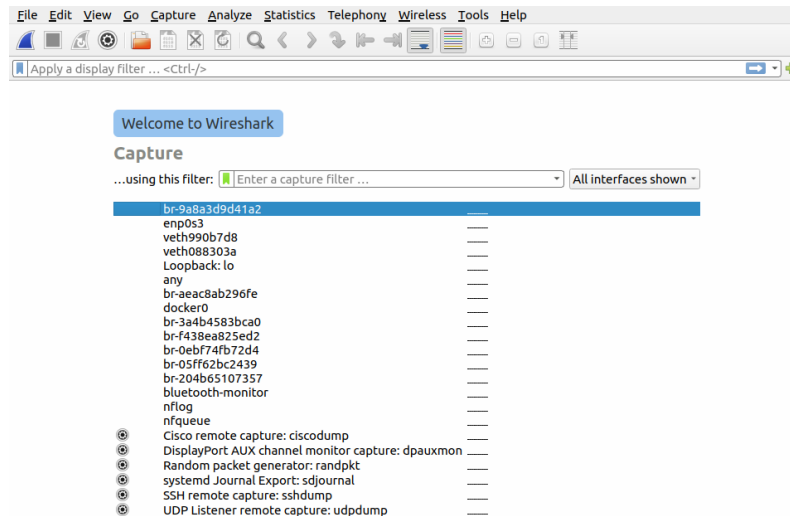
```
stephaniesalgado@VM: ~/Share/Labset... x
root@VM:/volumes# nano task1.1.py
root@VM:/volumes# ./task1.1.py
SNIFFING PACKETS.....

===== PACKET: 1 =====

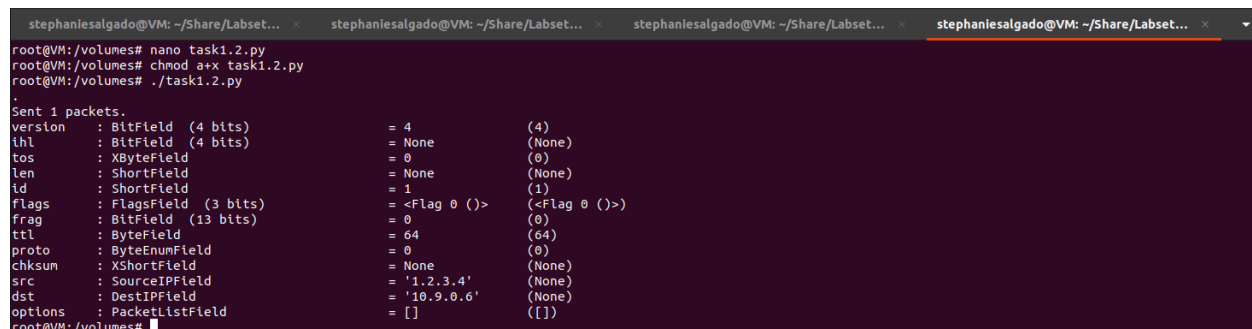
###[ Ethernet ]###
dst      = 02:42:71:7e:fa:24
src      = 02:42:0a:09:00:06
type     = IPv4
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x0
len      = 84
id       = 15396
flags    = DF
frag     = 0
ttl      = 64
proto    = icmp
chksum   = 0x7385
src      = 10.9.0.6
dst      = 128.230.0.11
\options \
###[ ICMP ]###
type     = echo-request
code     = 0
chksum   = 0xfade
id       = 0x3
seq      = 0x1
###[ Raw ]###
```

This was the output after filtering with “filter='net 128.230.0.0/16'” for subnet.

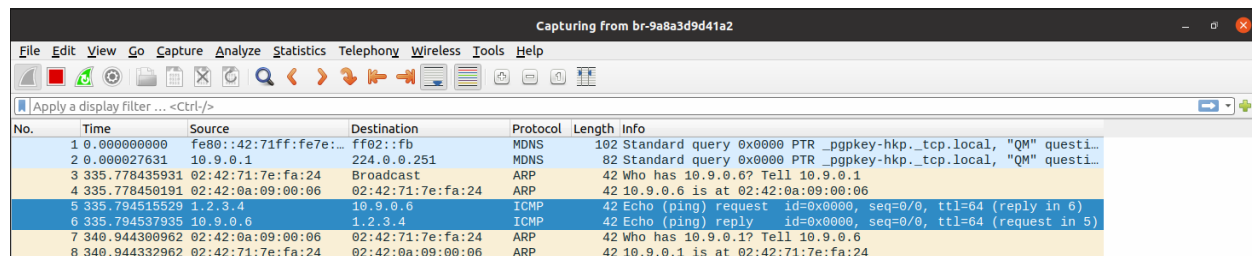
Task 1.2: Spoofing ICMP Packets



For this task, we will be using Wireshark to observe whether our request will be accepted by the receiver. I began by making sure I selected the correct interface (br-9a8a3d9d41a2).



I modified the provided task 1.2 code then made it executable and ran it.



Then I went to Wireshark to check if the request was accepted. I verified an echo reply packet was sent.

Task 1.3: Traceroute

```
stephaniesalgado@VM: ~/Share/Labset... x
root@VM:/volumes# touch task1.3.py
root@VM:/volumes# nano task1.3.py
root@VM:/volumes# chmod a+x task1.3.py
root@VM:/volumes# ./task1.3.py
1 10.0.2.2
2 192.168.1.1
3 70.94.136.1
4 24.28.133.233
5 24.175.34.16
6 24.175.32.149
7 24.175.32.156
8 131.150.63.17
9 108.170.225.147
10 142.251.253.25
11 8.8.4.4
root@VM:/volumes#
```

For this task, I asked for the destination IP, “8.8.4.4” which is the secondary DNS server for Google Public DNS. To accomplish this, I used the “sr1()” method from Scapy which listens and waits for a packet response. Eleven routers were reached before making it to the IP address destination.

Task 1.4: Sniffing and-then Spoofing

```
stephaniesalgado@VM: ~/Share/Labset... x
root@VM:/volumes# nano task1.4.py
root@VM:/volumes# chmod a+x task1.4.py
root@VM:/volumes# ./task1.4.py
Original Packet.....
Source IP: 10.9.0.6
Destination IP: 1.2.3.4
Spoofed Packet.....
Source IP: 1.2.3.4
Destination IP: 10.9.0.6
Spoofed Packet.....
Source IP: 10.9.0.6
Destination IP: 1.2.3.4
Original Packet.....
Source IP: 10.9.0.6
Destination IP: 1.2.3.4
```

I start editing the “task1.4.py” file, then making it executable and running it. I leave it running then go to “hostB”.

```
stephaniesalgado@VM: ~/Share/Labset... x stephaniesalgado@VM: ~/S
root@f589e678e850:/# ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
64 bytes from 1.2.3.4: icmp_seq=1 ttl=64 time=51.0 ms
64 bytes from 1.2.3.4: icmp_seq=2 ttl=64 time=155 ms
64 bytes from 1.2.3.4: icmp_seq=3 ttl=64 time=16.4 ms
64 bytes from 1.2.3.4: icmp_seq=4 ttl=64 time=14.5 ms
^C
--- 1.2.3.4 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3012ms
rtt min/avg/max/mdev = 14.456/59.163/154.780/57.089 ms
```

From “hostB”, I “ping 1.2.3.4” which should be a non-existing host on the Internet. Packets were still transmitted/received. As can be seen on the other screenshot, the output was simply flipped IPs from “Original Packet” to “Spoofed Packet”.

```
stephaniesalgado@VM: ~/Share/Labset... x stephaniesalgado@VM: ~/Share/Labset... x
stephaniesalgado@VM:~/Share/Labsetup$ sudo docker exec -it hostA-10.9.0.5 bash
root@2364903ac6b4:/# ping 10.9.0.99
PING 10.9.0.99 (10.9.0.99) 56(84) bytes of data.
From 10.9.0.5 icmp_seq=1 Destination Host Unreachable
From 10.9.0.5 icmp_seq=2 Destination Host Unreachable
From 10.9.0.5 icmp_seq=3 Destination Host Unreachable
^C
--- 10.9.0.99 ping statistics ---
5 packets transmitted, 0 received, +3 errors, 100% packet loss, time 4126ms
pipe 3
```

I used “hostA” this time. Notice that there was “100% packet loss”.

```
stephaniesalgado@VM: ~/Share/Labset... x
root@VM:/volumes# ./task1.4.py
^Croot@VM:/volumes#
```

The output was blank. This is because since the host does not exist on LAN, the packet can’t even be created.


```
stephaniesalgado@VM: ~/Share/Labset... x
root@VM:/volumes# ./task1.4.py
Original Packet.....
Source IP: 10.9.0.5
Destination IP: 8.8.8.8
Spoofed Packet.....
Source IP: 8.8.8.8
Destination IP: 10.9.0.5
Spoofed Packet.....
Source IP: 10.9.0.5
Destination IP: 8.8.8.8
Spoofed Packet.....
Source IP: 10.9.0.5
Destination IP: 8.8.8.8
```

Once again, after modifying the file, I ran it and then went to “hostA” to “ping 8.8.8.8”, which is the existing host on the internet. The output seemed similar to the output I got for the non-existing host on the internet, in which the IPs on the original and spoofed packets switched.

```
stephaniesalgado@VM: ~/Share/Labset... x stephaniesalgado@VM: ~/Share/Labset... x
root@2364903ac6b4:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=113 time=33.0 ms
64 bytes from 8.8.8.8: icmp_seq=1 ttl=64 time=40.0 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=2 ttl=64 time=18.5 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=113 time=35.0 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=3 ttl=64 time=17.3 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=113 time=38.2 ms (DUP!)
^C
--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, +3 duplicates, 0% packet loss, time 2007ms
rtt min/avg/max/mdev = 17.277/30.322/39.993/9.083 ms
```

Unlike when pinging “1.2.3.4”, which is the host that doesn’t exist on the internet, pinging “8.8.8.8” shows “DUP !” (duplicates). This basically serves as a response from the real host.

Summary:

Once again, this lab made good use of containers. I had never used Wireshark before so that's at least one take away from the lab. Scapy seems like a very valuable tool when working on things like this. I found it interesting how you can make a program and use the library to filter for specific sources and destinations. I was surprised by the fact that a library such as Scapy can be used to implement your own versions of traceroute. Furthermore, being able to tell whether a host is existing on the internet or LAN based off of the output from Task 1.4, means that using similar sniffing and spoofing programs can be incredibly valuable to attackers.